

Experimental Evaluation of a Multi-Layer Abstraction Approach for RL on Non-Markovian Tasks

Advisor: Prof. Fabio Patrizi

Master of Science in Engineering in Computer Science

Matteo Ventali

Sapienza University of Rome



SAPIENZA
UNIVERSITÀ DI ROMA



Outline

- ▶ Motivation
- ▶ Framework and Implementation
- ▶ Experimental Evaluation



Motivation and Problem Statement

1 Motivation

Sparse Rewards

Informative feedback
is received only after
reaching the final goal.

Non-Markovian Task


Task completion may depend
on the history of previously
observed events.

Prohibitive Learning



The Starting Point

2 Framework and Implementation




Framework by Cipollone et al. [1]



Extension Proposed in This Thesis

2 Framework and Implementation






Experimental Setting

3 Experimental Evaluation

Ground Domain M_0

Gymnasium Lunar Lander




- $S = (x, y, v_x, v_y, \theta, \omega, c_l, c_r)$
- $\mathcal{A} = \{\text{no-op, eng-left, eng-right, eng-main}\}$
- Original dense reward removed

State mapping



Abstract Domain M_i

Grid Environment




- $\mathcal{S}_i = \{(i,j) : 1 \leq i \leq H, 1 \leq j \leq W\}$
- $\mathcal{A} = \{\uparrow, \nearrow, \rightarrow, \searrow, \downarrow, \swarrow, \leftarrow, \nwarrow\}$
- Physical quantities discarded




Biased-Unbiased Interaction: Practical Limitation

3 Experimental Evaluation

Episodic Tasks







1-phase: $\diamond g_1$



2-phase: $\diamond(wp_1 \wedge \bigcirc \diamond g_1)$

Training results



Greedy eval.: standalone unbiased (left), coupled biased/unbiased (right); phases 1/2 (top/bottom).



Reward Shaping Under Temporal Misalignment


3 Experimental Evaluation

Shaping Formulation

$$\begin{cases} F(s, s') = \gamma_s \Phi(s') - \Phi(s) \\ \Phi(s) = V(s) \end{cases}$$

Potential Based Reward Shaping (PBRs) [1, 2].

Training Results



γ_s	Best success rate (%)	Same-state fraction (%)	Same-state reward per transition	Same-state reward per episode
0.99	0.00 $\pm$ 0.00	82.39 $\pm$ 0.96	-89.51 $\pm$ 0.05	-5000.60 $\pm$ 53.21
1.00	96.94 $\pm$ 1.91	91.48 $\pm$ 0.22	0.00 $\pm$ 0.00	0.00 $\pm$ 0.00

Training reward (left) and greedy evaluation (right), with transition statistics for $\gamma_s = 0.99$ and $\gamma_s = 1$.



From Theory to Practice

3 Experimental Evaluation

Theory

Coupled architecture

Biased + Unbiased

PBRs Formulation

$$F(s, s') = \gamma_s V(s') - V(s)$$

Practice

Biased Learner

Only biased on the ground MDP

Shaping Heuristic

$$F(s, s') = V(s') - V(s)$$

Interaction
Problem


Temporal
Mismatch



Scaling to More Complex Non-Markovian Tasks

3 Experimental Evaluation

Required Tasks



Training Results

Task	Baseline	6×6	12×12	24×24	48×48	180×180	360×360
2-phase	16.1	91.9	97.4	96.8	98.3	95.1	97.0
3-phase Simple	5.3	95.0	94.1	98.2	98.3	98.8	96.5
3-phase Hard	0.1	95.0	93.1	97.0	88.0	98.6	97.1
3-phase Extreme	0.0	28.3	77.8	93.2	87.6	90.1	91.6
4-phase	0.0	24.8	53.5	83.6	74.9	64.2	85.1
5-phase	0.0	52.3	62.3	65.2	87.3	88.5	85.6

Best-policy success rate (%), mean over 10 training seeds.



Learning Reliability

3 Experimental Evaluation



Reliability vs. Quality

Grid	4-phase $N \geq 80\%$	4-phase $SR \geq 80\%$	5-phase $N \geq 80\%$	5-phase $SR \geq 80\%$
6×6	1/10	84.2%	4/10	94.6%
12×12	3/10	89.9%	5/10	97.5%
24×24	6/10	94.3%	6/10	95.9%
48×48	8/10	93.4%	9/10	96.9%
180×180	7/10	91.2%	9/10	98.4%
360×360	9/10	94.8%	9/10	95.1%

$N \geq 80\%$: number of high-performing seeds.

Mean $SR \geq 80\%$: mean post-training best-policy success rate among high-performing seeds.

Seed-Level Distributions




Top: 4-phase trajectory; bottom: 5-phase trajectory.
Each point represents one training seed (10 seeds per configuration).



Learning on Abstract Levels


3 Experimental Evaluation

Learning Stacks



$90 \times 90 \rightarrow 360 \times 360 \rightarrow \mathcal{M}_0$ vs $360 \times 360 \rightarrow \mathcal{M}_0$ (right).

Training Results




Task success rate during training under uniform and Gym-reset evaluation.



From a Learned Policy to an Effective Shaping Potential


3 Experimental Evaluation

Value-Function Heatmaps



Value functions for VI, policy evaluation, and MaxQ (top: q_1 ; bottom: q_2).

Training Results




Task reward and episode success during training.




Learning on Cyclic Tasks

3 Experimental Evaluation


Cyclic Tasks




Horizontal



Vertical




Triangle



Rectangle

Training Results





Evaluation episodes grouped by completed cycles (0-3), for baseline and shaped policies.



Additional Case Studies



3 Experimental Evaluation

Task with Constraints



Shaped reward rises, while the baseline remains near zero.

Landing Task



Landing reward improves during training, with variation across seeds.



Conclusions, Limitations & Future Work

3 Experimental Evaluation

Conclusions

- Effective abstraction-based guidance
- Automatic guidance generation from logical formulas
- Multi-level architecture is promising

Limitations

- Coupled architecture interaction
- Temporal misalignment between levels
- Abstract policies could be not effective as guidance
- Long temporal tasks are difficult to be discovered

Future Works

- Study new exploration mechanism
- Introduce different techniques across levels
- Richer logical formalism



Thank you for listening!
Any questions?



References

3 Experimental Evaluation

- [1] Roberto Cipollone et al. "Exploiting robot abstractions in episodic RL via reward shaping and heuristics". In: *Robotics and Autonomous Systems* 193 (2025), p. 105116. DOI: 10.1016/j.robot.2025.105116.
- [2] Andrew Y. Ng, Daishi Harada, and Stuart J. Russell. "Policy Invariance Under Reward Transformations: Theory and Application to Reward Shaping". In: *Proceedings of the Sixteenth International Conference on Machine Learning*. ICML '99. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 1999, 278–287. ISBN: 1558606122.
- [3] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. Second. The MIT Press, 2018.
- [4] Volodymyr Mnih et al. "Human-level control through deep reinforcement learning". In: *Nature* 518.7540 (2015), pp. 529–533. DOI: 10.1038/nature14236.
- [5] Hado van Hasselt, Arthur Guez, and David Silver. "Deep Reinforcement Learning with Double Q-Learning". In: *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*. 2016, pp. 2094–2100.
- [6] Christopher J. C. H. Watkins and Peter Dayan. "Q-Learning". In: *Machine Learning* 8 (1992), pp. 279–292. DOI: 10.1007/BF00992698.
- [7] Giuseppe De Giacomo and Moshe Vardi. "Linear temporal logic and Linear Dynamic Logic on finite traces". In: *IJCAI International Joint Conference on Artificial Intelligence* (Jan. 2013), pp. 854–860.
- [8] Amir Pnueli. "The temporal logic of programs". In: *18th Annual Symposium on Foundations of Computer Science (sfcs 1977)*. 1977, pp. 46–57. DOI: 10.1109/SFCS.1977.32.
- [9] Giorgio Ausiello, Fabrizio D'Amore, and Giorgio Gambosi. *Linguaggi, modelli, complessità*. Franco Angeli, 2003. ISBN: 9788846444707.
- [10] Fahiem Bacchus, Craig Boutilier, and Adam Grove. "Rewarding behaviors". In: *Proceedings of the Thirteenth National Conference on Artificial Intelligence*. AAAI'96. Portland, Oregon: AAAI Press, 1996, 1160–1167.
- [11] Ronen I. Brafman, Giuseppe De Giacomo, and Fabio Patrizi. "LTLf/LDLf Non-Markovian Rewards". In: *AAAI Conference on Artificial Intelligence*. 2018.
- [12] Rodrigo Toro Icarte et al. "Reward Machines: Exploiting Reward Function Structure in Reinforcement Learning". In: *J. Artif. Int. Res.* 73 (May 2022). ISSN: 1076-9757. DOI: 10.1613/jair.1.12440. URL: <https://doi.org/10.1613/jair.1.12440>.
- [13] Giuseppe De Giacomo et al. "Restraining Bolts for Reinforcement Learning Agents". In: *The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, The Thirty-Second Innovative Applications of Artificial Intelligence Conference, IAAI 2020, The Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020*, New York, NY, USA, February 7–12, 2020. AAAI Press, 2020, pp. 13659–13662. DOI: 10.1609/AAAI.V34I09.7114. URL: <https://doi.org/10.1609/aaai.v34i09.7114>.
- [14] Mark Towers et al. *Gymnasium: A Standard Interface for Reinforcement Learning Environments*. 2025. arXiv: 2407.17032 [cs.LG]. URL: <https://arxiv.org/abs/2407.17032>.
- [15] Farama Foundation. *Lunar Lander*. URL: https://gymnasium.farama.org/environments/box2d/lunar_lander/ (visited on 10/08/2026).
- [16] Francesco Fuggitti. *LTLf2DFA*. Version 1.0.3. 2019. DOI: 10.5281/zenodo.3888410.
- [17] Marek Grzes. "Reward Shaping in Episodic Reinforcement Learning". In: *Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems*. AAMAS '17. São Paulo, Brazil: International Foundation for Autonomous Agents and Multiagent Systems, 2017, 565–573.
- [18] Ziyu Wang et al. *Dueling Network Architectures for Deep Reinforcement Learning*. 2016. arXiv: 1511.06581 [cs.LG]. URL: <https://arxiv.org/abs/1511.06581>.
- [19] Diederik P. Kingma and Jimmy Ba. *Adam: A Method for Stochastic Optimization*. 2017. arXiv: 1412.6980 [cs.LG]. URL: <https://arxiv.org/abs/1412.6980>.